

Optimizing Disk Defragmentation Using Quantum Computation

Raghav Yadav, Maulik Gupta, Parth Singh, Rishank Mishra, Darshita Yadav

BML Munjal University 67th Milestone, NH-48, Sidhrawali, Gurugram, Haryana – 122413, India.

Abstract

Filesystem fragmentation remains a primary cause of performance degradation in modern computer systems[1]. As files are modified over time, data segments become scattered, severely degrading Input/Output (I/O) throughput, increasing latency, and accelerating mechanical wear in data centers. Current defragmentation and allocation strategies rely on classical heuristic algorithms, such as First-Fit and Best-Fit, to manage file blocks. However, these classical methods face fundamental scalability barriers; specifically, the "Best-Fit" strategy requires a sequential search of free storage blocks, resulting in linear time complexity ($O(N)$). As storage volumes expand into the multi-terabyte regime, this linear scan creates a significant I/O bottleneck. To address this limitation, this research proposes a quantum-enhanced framework for disk allocation using Grover's Algorithm[2]. Unlike classical sequential searches, our approach leverages quantum superposition to query the entire address space of the disk simultaneously. By utilizing Amplitude Amplification, the proposed system identifies the optimal "Best-Fit" storage block with a quadratic speedup ($O(\sqrt{N})$). This paper details the circuit construction for this quantum search protocol and evaluates its potential to mitigate the latency constraints of classical storage management.

Introduction

Research into the optimization of mechanical storage devices has been a central pillar of computer science for over half a century. The fundamental algorithms governing disk head scheduling were established in the late 1960s, most notably with Denning's introduction of the SCAN algorithm in 1967 [3].

These early theoretical foundations laid the groundwork for filesystem designs such as the Berkeley Fast File System (FFS) in 1984, which explicitly attempted to mitigate fragmentation through intelligent block placement policies[4].

Despite these decades of algorithmic refinement, filesystem fragmentation remains a persistent bottleneck in the era of big data and hyperscale computing. As files are created, modified, and deleted over time, valid data segments become scattered across non-contiguous blocks of the

storage medium. This physical separation forces storage controllers to perform excessive seek operations, leading to severe degradation in Input/Output (I/O) throughput, increased latency, and accelerated mechanical wear. Furthermore, in large-scale data centers, these inefficiencies compound, resulting in substantial energy waste and reduced hardware lifecycles.

To overcome these classical limitations, this research proposes a novel, quantum-enhanced framework for disk allocation. We posit that Quantum Computing offers a specific advantage in unstructured search problems through the application of Grover's Algorithm[2]. By mapping the free-space bitmap of a drive to a quantum register, we can translate the physical problem of finding a free block into a database search problem solvable by quantum processors. By leveraging quantum mechanical phenomena such as superposition and entanglement, our proposed framework can evaluate the suitability of all storage blocks simultaneously. This approach addresses the scalability gap of classical heuristics, offering a theoretical quadratic reduction in the time required to calculate optimal write paths. This paper details the mapping of filesystem logic to quantum circuits and evaluates the feasibility of this quantum advantage in next-generation storage systems.

Keywords

Quantum Computing, Disk Allocation, Grover's Algorithm, Storage Optimization, Amplitude Amplification, File Systems.

Classical Defragmentation

Most existing disk defragmentation strategies depend on classical heuristics, such as First-Fit and Best-Fit, which prioritize immediate local optimization over global efficiency. While these "greedy" algorithms are computationally cheap in the short term, they are mathematically suboptimal for long-term storage health.


For instance, by consistently attempting to fill the smallest available gaps, the Best-Fit strategy inadvertently fractures free space into unusable micro-segments, effectively maximizing external fragmentation. Furthermore, because these methods rely on sequential searches, their complexity scales linearly $O(N)$.

This creates a significant Input/Output (I/O) bottleneck as storage capacities expand into the multi-terabyte range.

$$j^* = \operatorname{argmin}_{i \in \mathcal{F}, C(h_i) \geq S_{req}} (C(h_i) - S_{req})$$

Equation (1)

The selection logic for the Best Fit strategy is expressed in Equation (1). By solving for the index j^* that minimizes the residual capacity $(C(h_i) - S_{\{req\}})$, the algorithm attempts to maximize storage density for individual write events. However, as demonstrated in the comprehensive review by Wilson et al. [5], this greedy minimization of internal fragmentation frequently exacerbates external fragmentation, leading to a 'checkerboard' of unusable micro-segments that necessitate the global defragmentation solution we propose in Section III.



Checkerboard Pattern in Fragmented Disk
(generated using gParted in Ubuntu 24LTS)

Quantum Optimization

We propose implementing Grover's Algorithm, a quantum search method that fundamentally changes the rules of the game. Instead of checking each storage block individually $O(N)$, Grover's algorithm allows the storage controller to query the probability of all blocks at once, amplifying the signal of the "Best Fit" candidate while suppressing the others. This provides a quadratic speedup $O(\sqrt{N})$, effectively allowing us to find the optimal write location in a fraction of the time required by classical methods.

Motivation for the Quantum Approach

The primary motivation for transitioning to a quantum-enhanced framework lies in the inherent scalability limitations of classical computing architectures. As global data storage needs expand into the petabyte and exabyte regimes, the linear time complexity ($O(N)$) of traditional allocation heuristics—specifically the sequential scan required for the 'Best-Fit' strategy—becomes a prohibitive bottleneck. In modern hyperscale data centers, this linear dependency means that every increase in storage capacity directly correlates to an increase in Input/Output (I/O) latency,

effectively punishing the system for growing larger. The quantum approach circumvents this serial constraint by leveraging superposition to evaluate the suitability of all available storage blocks simultaneously rather than sequentially. By utilizing Grover's Algorithm to achieve a quadratic speedup ($O(N)$), we offer a computational paradigm that significantly decouples allocation speed from total disk capacity, ensuring that system performance remains robust even as data volumes scale exponentially.

Circuit Construction and Gate Logic

To implement this search on a gate-based quantum processor, we construct a circuit that iterates through two main phases. The qubits in this circuit do not store the data itself, but rather the index addresses of the free blocks.

The Hadamard gate (H) serves as the fundamental catalyst for quantum parallelism within the proposed circuit architecture. Functioning as a single-qubit unitary operator, it maps the deterministic computational basis states, $|0\rangle$ and $|1\rangle$, into coherent states of equal-weight superposition (denoted as $|+\rangle$ and $|-\rangle$). In the context of the disk defragmentation search, the first stage of the algorithm applies a tensor product of Hadamard gates $H^{\{\otimes n\}}$ to the initialized address register $|0\rangle^{\{\otimes n\}}$. This operation transforms the system from a single null state into a uniform superposition of all 2^n possible storage block indices. This step is mathematically prerequisite to the search process; by spreading the probability amplitude equally across the entire configuration space, it allows the subsequent Oracle operator to evaluate the "Best Fit" suitability of every physical disk sector simultaneously, rather than sequentially.

1. Initialization: Creating the Superposition (The "H" Layer)

The circuit begins with all qubits initialized to the zero state. We then apply a layer of Hadamard Gates (H) to every qubit.

- In a classical computer, 3 bits can only represent *one* number at a time (001, 010...). The Hadamard gates force the qubits into a state of "Uniform Superposition."
- The system now simultaneously represents every possible address on the hard drive with equal probability. This is the equivalent of looking at the entire haystack at once, rather than picking up one piece of straw at a time.

2. The Oracle U_f :


The system then passes through the Oracle block. This is a black-box function that contains the logic for our "Best Fit" criteria (e.g., "Is this block > 100KB?").

- The Oracle checks every address in the superposition simultaneously.
- If an address is *not* the correct one, the Oracle does nothing. If the address *is* the Best Fit, the Oracle flips its phase (multiplying the state by -1).

3. The Diffuser U_s :

Amplifying the Signal, the diffuser acts by making the right block more prominent in the result.

- The Diffuser acts as a mathematical mirror. It calculates the average probability amplitude of all states and "flips" them around that average.
- Because the right block was "upside down" (negative phase), this flip causes its probability to skyrocket (constructive interference), while the probabilities of all the wrong answers cancel out and shrink (destructive interference).



Circuit diagram of the quantum setup

Methods:

Classical Implimentation in Python

References

- [1] J. Park and Y. I. Eom, 'Filesystem Fragmentation on Modern Storage Systems', *ACM Trans. Comput. Syst.*, vol. 41, no. 1–4, Dec. 2023, doi: 10.1145/3611386.

- [2] L. K. Grover, 'A fast quantum mechanical algorithm for database search', in *Proceedings of the twenty-eighth annual ACM symposium on Theory of computing*, 1996, pp. 212–219.
- [3] P. J. Denning, 'Effects of scheduling on file memory operations', in *Proceedings of the April 18-20, 1967, Spring Joint Computer Conference*, in AFIPS '67 (Spring). New York, NY, USA: Association for Computing Machinery, 1967, pp. 9–21. doi: 10.1145/1465482.1465485.
- [4] M. K. McKusick, W. N. Joy, S. J. Leffler, and R. S. Fabry, 'A fast file system for UNIX', *ACM Transactions on Computer Systems (TOCS)*, vol. 2, no. 3, pp. 181–197, 1984.
- [5] P. R. Wilson, M. S. Johnstone, M. Neely, and D. Boles, 'Dynamic storage allocation: A survey and critical review', in *International Workshop on Memory Management*, Springer, 1995, pp. 1–116.